



# Paging


Mythili Vutukuru  
CSE, IIT Bombay

# Recap: Paging and page table




# Page table entry

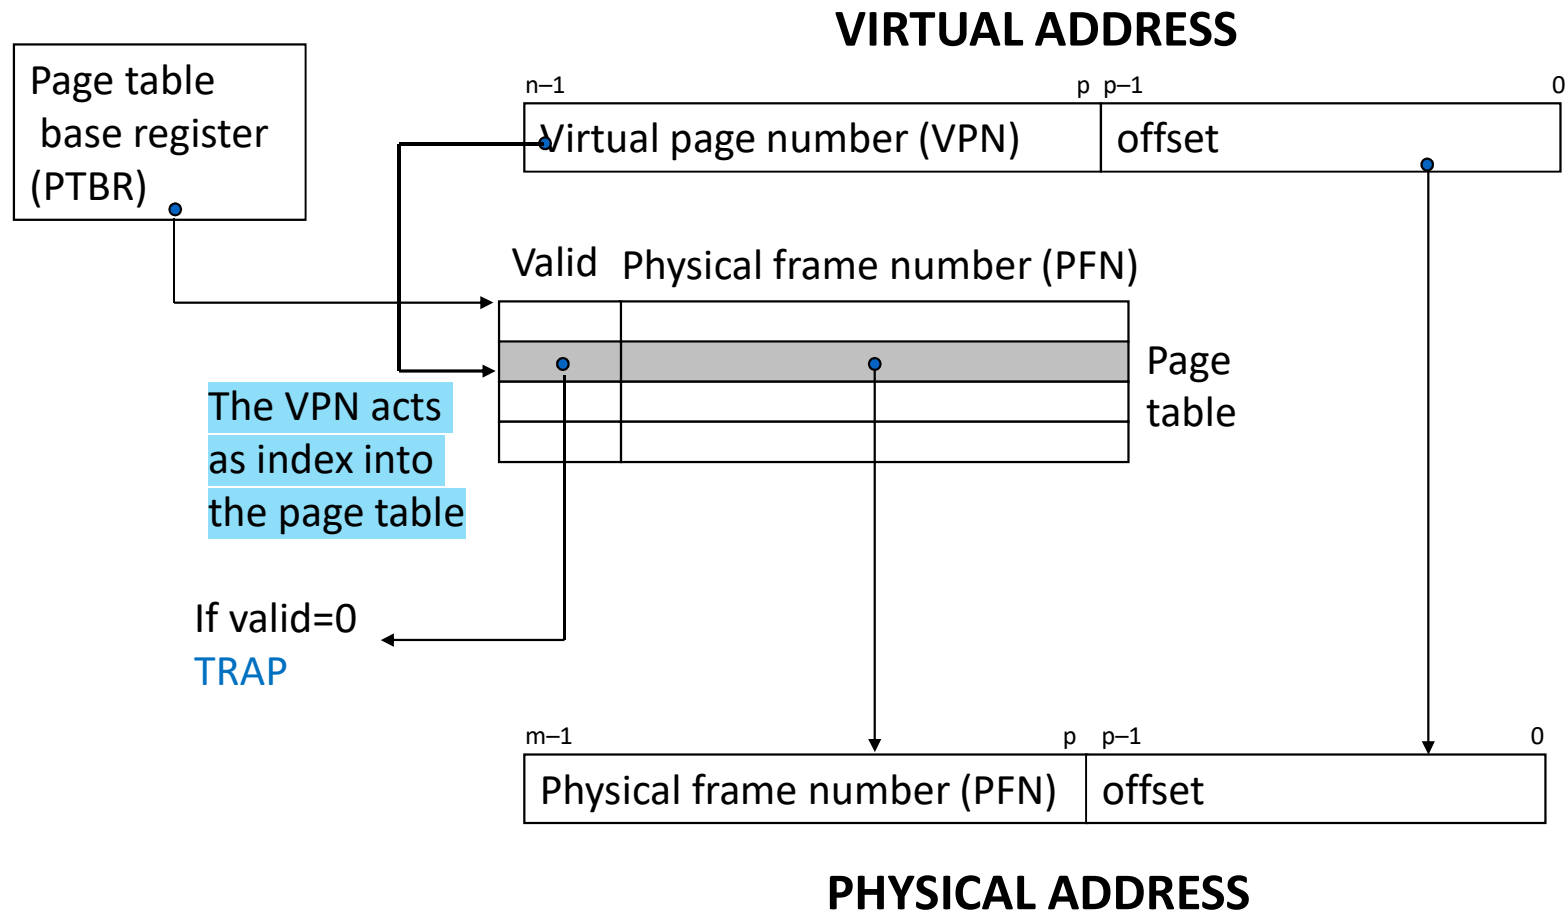
- Page table is array of page table entries, one per page of process
- i-th **page table entry (PTE)** contains physical frame number and other details (permissions, status, ..) of i-th page of process
  - Valid: is this page in use by process (not all virtual addresses are used by process)
  - Various permission bits (more later)
  - Other status bits: present, dirty, accessed (more later)



# Address translation in MMU

- MMU stores starting (physical) address of page table array in CPU register called page table base register
- Page size determines number of bits in offset
  - 4KB pages need  $\log_2(4K) = 12$  bits as offset within page
- Remaining most significant bits give VPN
  - For 32-bit machines and 4KB pages, 20 bit VPN
- MMU uses VPN as index into page table array, accesses PTE, gets PFN, adds offset bits to get PA
- If no valid PTE found, MMU traps to OS





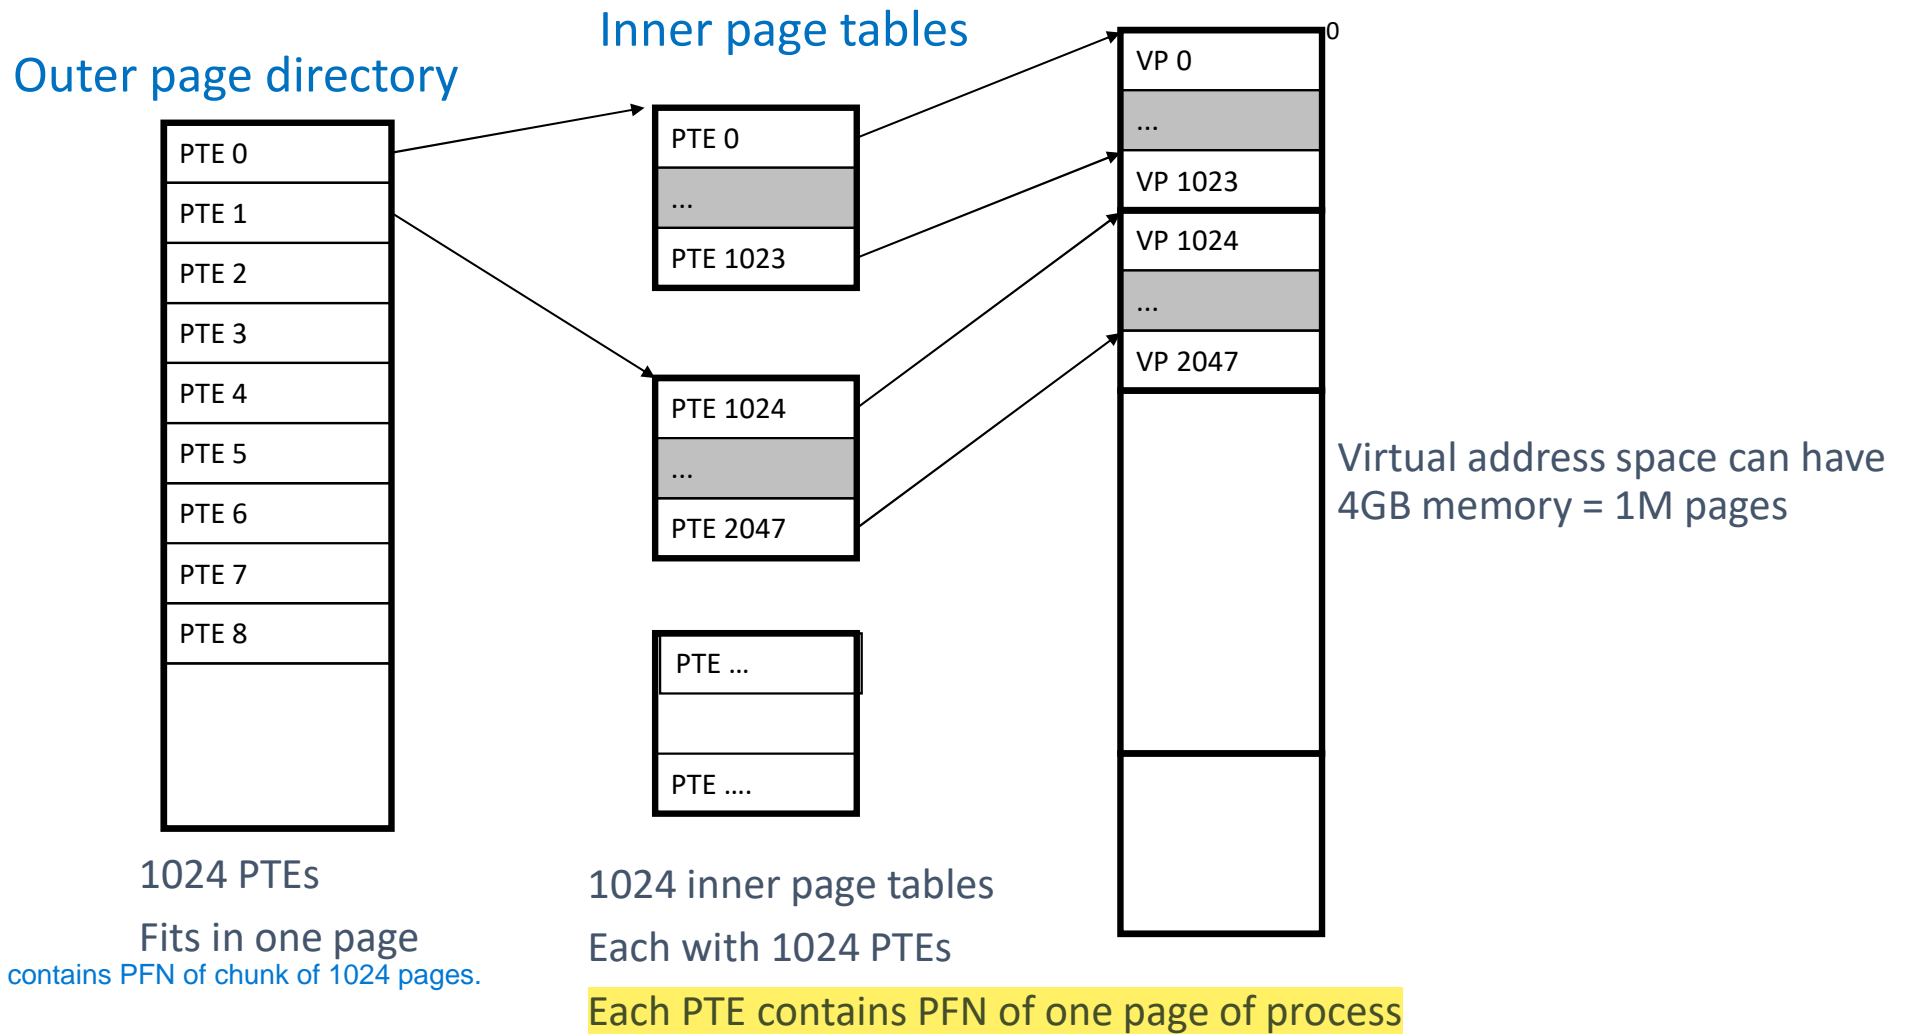
# Size of page tables

$1K = 2^{10} = 1024$   
 $1M = 2^{20} = 1024 * 1024$   
 $1G = 2^{30} = 1024 * 1024 * 1024$   
B = byte, b = bit

- What is typical size of page table in a 32-bit system?
- $2^{32} = 4GB$  virtual address space
- Assume page size =  $4KB = 2^{12}$
- Number of PTEs = number of pages in virtual address space =  $(2^{32}/2^{12}) = 2^{20} = 1M$
- If each PTE is 4 bytes, page table size = 4 bytes \* 1M entries = 4MB
- How are page tables stored in memory?
  - All memory is only allocated in 4KB chunks, so how to store 4MB?
- Solution: split page table into pages (much like memory image), use another page table to keep track of original page table!

# Two-level page table in 32-bit systems

- 4MB page table split into 1024 chunks of 4KB each (to fit in page)
- 1M PTEs split across 1024 pages, each containing 1024 PTEs
- Physical frame numbers of these 1024 chunks stored in an outer page table or page directory
  - 4 byte page table entry each, so outer page directory fits in one page here
- Page table now has two levels
  - **Outer page table (page directory)** has physical frame numbers of 1024 “inner” page table pages
  - Each **inner page table** has physical frame numbers (PTEs) of 1024 pages of the process virtual address space

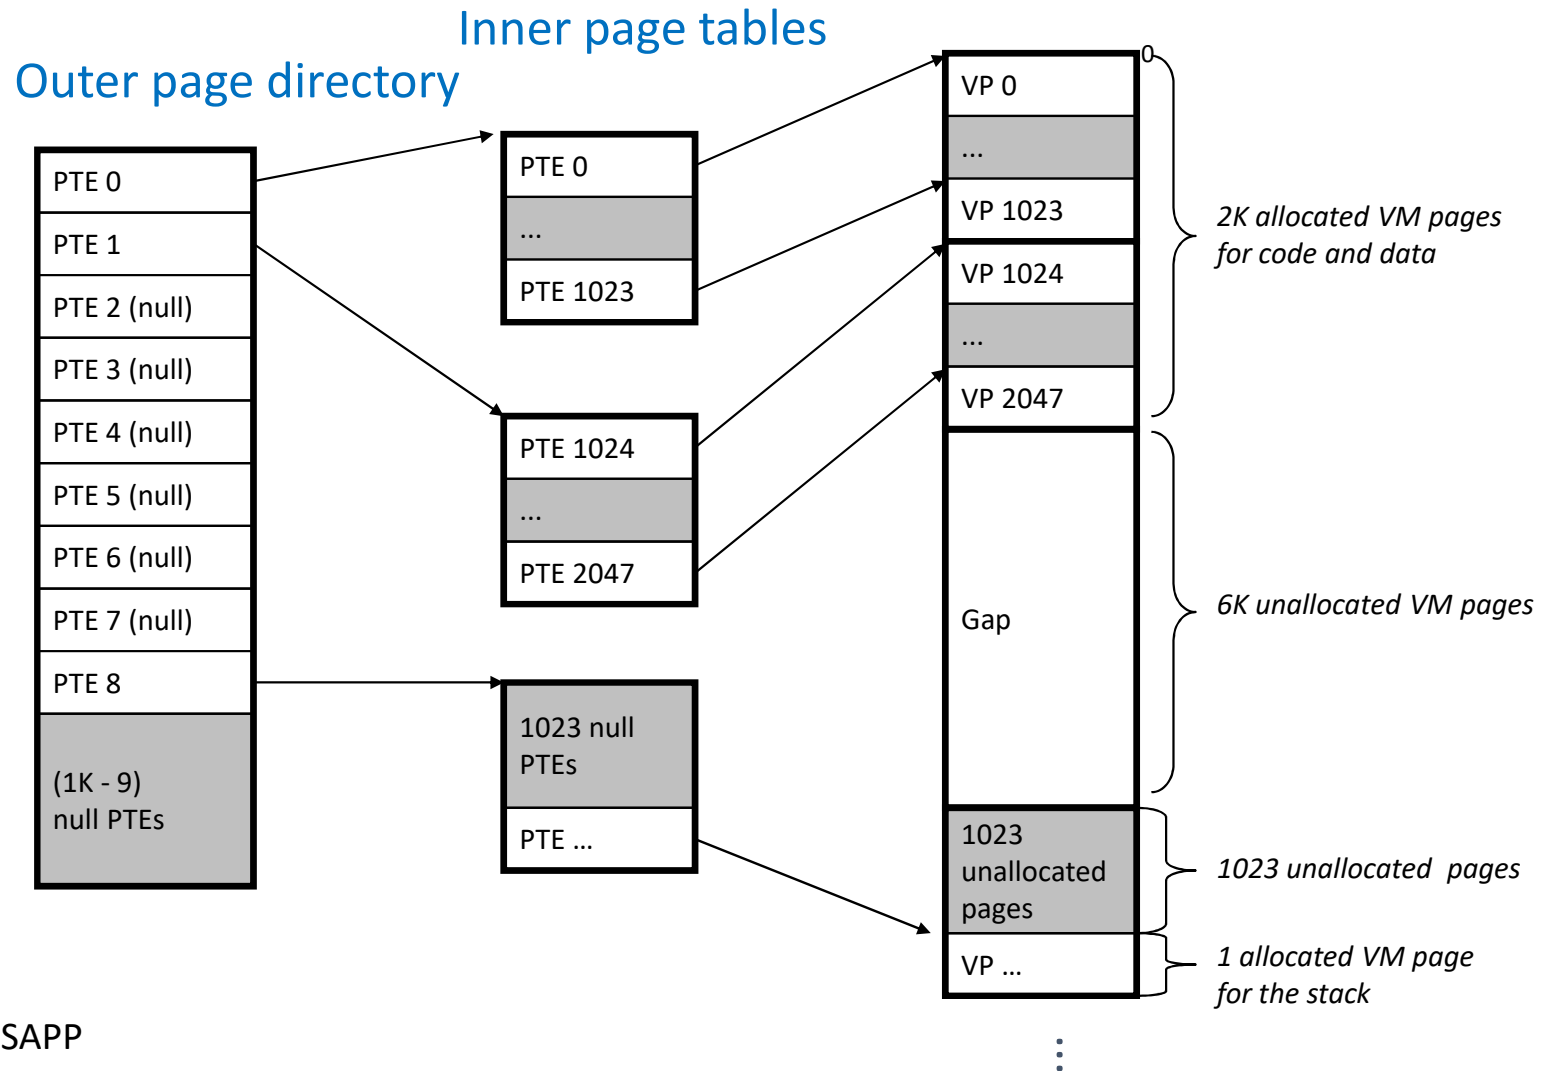




# Inner page tables on demand

- Note: not all inner page tables need to be created always, only those with at least one valid entry needed
- Example: Process with 2K pages of code+data, 6K + 1023 unallocated pages in address space, then one page allocated for stack
  - First two inner page tables are allocated, hold the 2K valid PTEs
  - Next 6 inner page tables are not created, the corresponding entries in outer page directory are invalid / null
  - In next inner page table, 1023 invalid entries and one valid PTE containing frame number of stack page
  - Remaining inner page tables not created, corresponding outer page directory entries are invalid

Array sizes will be fixed but the corresponding entries will not be created.



all the below entries won't be created at all.

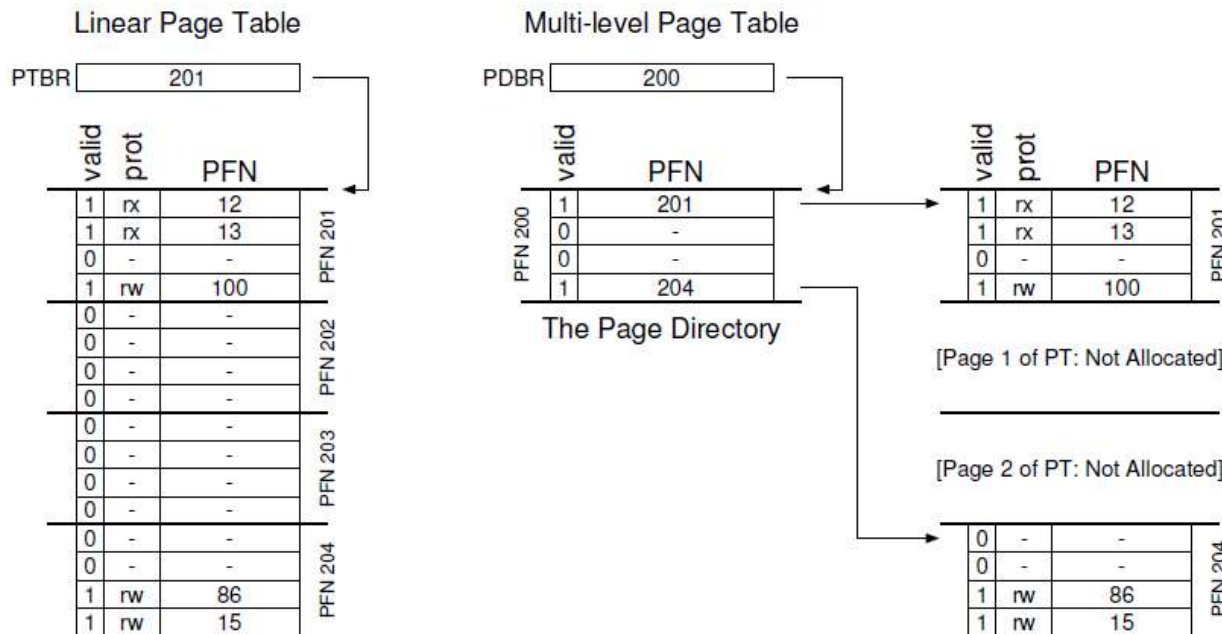
Outer 10 bits will have the address of physical frame which contains the address of fixed chunk of page table entries in it, then the next 10 bits will locate the actual page table entry within that and mmu will translate it into physical frame address and then using offset we'll get the content requested

# Address translation in 2-level page table

- Virtual address of 32 bits = 20 bit page number + 12 bit offset
- 20 bits index into a single page table is now used as
  - Most significant 10 bits index into page directory, locate PTE of one of the 1024 inner page tables contain our desired address
  - Next 10 bits index into inner page table to locate PTE of page
- Locate PTE, computer physical address using frame number and 12-bit offset into page
- MMU “walks” the multiple levels of the page table to translate virtual addresses [That's why TLB is a good helper in this case.](#)

# Page table/directory base register

- Single level: MMU stores starting address of page table in page table base register
- Multi-level: MMU stores starting address of outer page directory in page directory base register (CR3 register in x86)



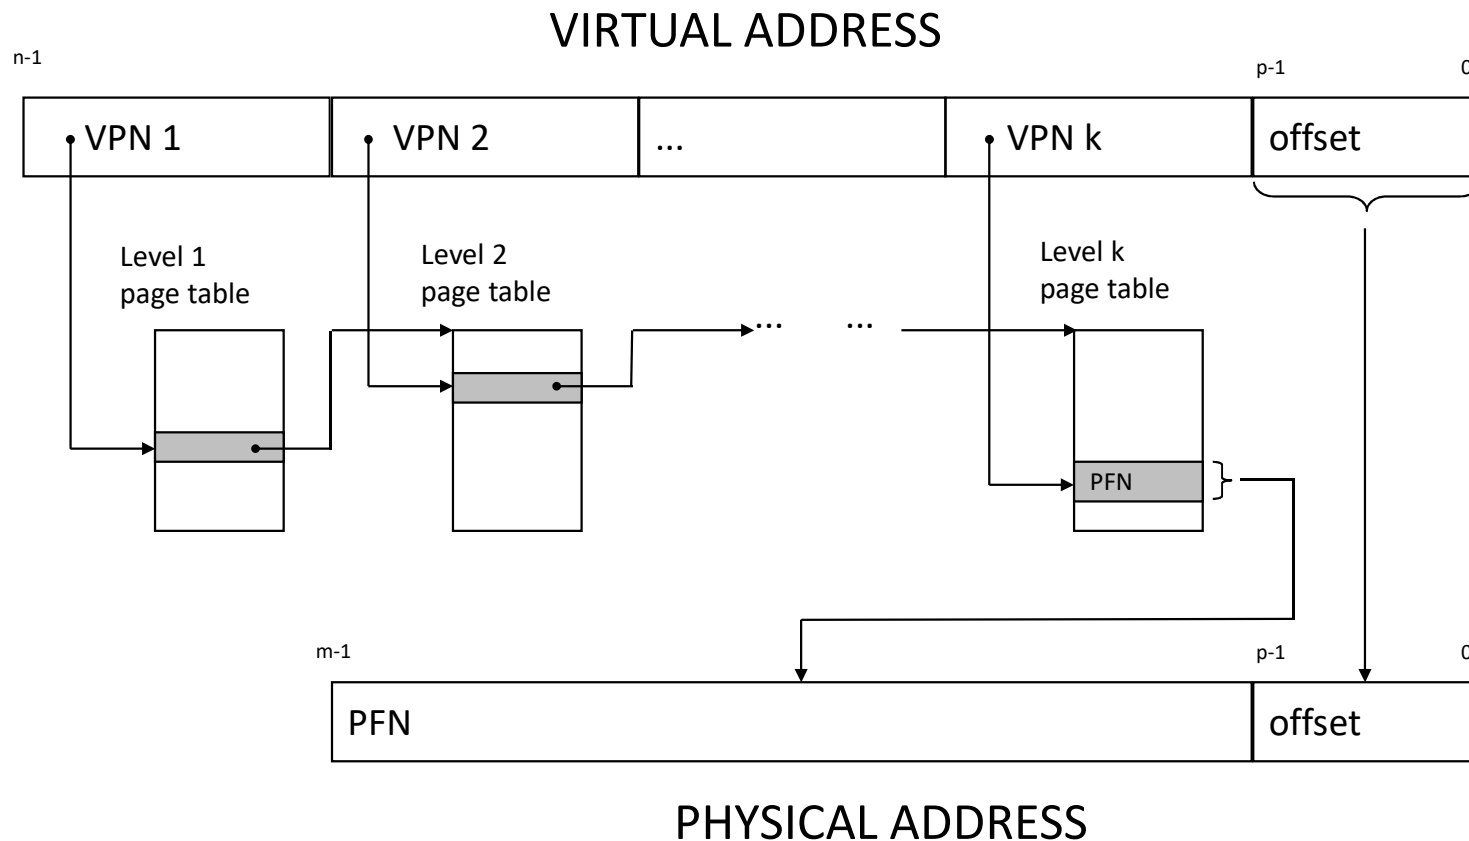
# Multi-level page tables

- What if outer page directory does not fit into one page?
- Store page directory across many pages, use yet another page table to store frame numbers of page directory pages
- This can go on until outermost page table fits in one page I need to work on this translation.
- Example: 48-bit CPU, 4KB pages, 8 byte page table entries
  - $2^{48}$  bytes in virtual address space =  $2^{36}$  pages for each process
  - Each page can store  $4\text{KB}/8 = 2^9 = 512$  page table entries
  - Innermost level (actual page table) has  $2^{36}$  page table entries = needs  $2^{27}$  pages
  - Innermost page table split into multiple pages =  $2^{27}$  page table entries to track innermost page table pages
  - Next level of page table stores  $2^{27}$  page table entries = needs  $2^{18}$  pages
  - Next level stores  $2^{18}$  page table entries = needs  $2^9 = 512$  pages
  - Outermost level can store all 512 page table entries in 1 page

# Address translation with 4-level page table

- Example: 48-bit CPU, 4KB pages, 8 byte page table entries
  - 4 level page table required
  - Outermost page directory has 512 entries, containing frame numbers of next level page table pages, each of those contain frame numbers of next level page table, ...
  - Page table at  $i$ -th level has frame numbers of 512  $(i+1)$ -th level page table pages
- How to translate VA to PA?
  - 48-bit VA = 36 bits + 12 bit offset
  - 36 bits = 9 bit offset into each of the 4 levels of page table
- If TLB miss, MMU has to access 4 different memory locations for 4 levels of page table, in order to translate one VA to PA
- MMU page table walks become even longer, TLB hit rate is critical

# Address translation with multi-level page table



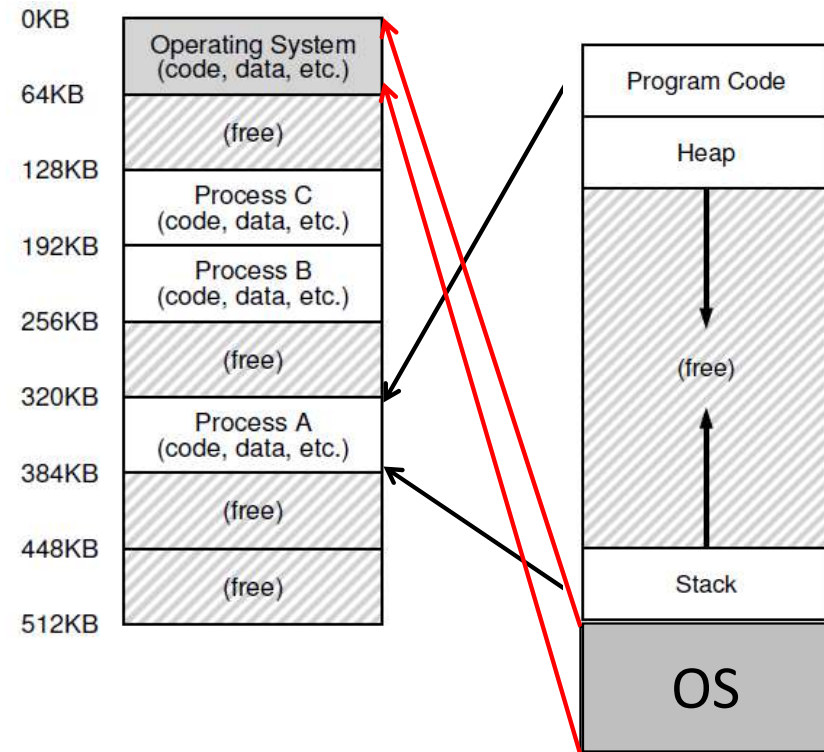
# Revisiting process virtual address space

- What should virtual address space/page table of process have? Any memory that the process needs to access during its execution
  - Its own memory image: code, data, stack, heap
  - Other common memory it needs to access: shared language libraries, OS
- Why? MMU allows access to memory **only via virtual addresses**
  - Can only access physical memory mapped in page table at some virtual address
  - So all physical memory needed by process should be mapped into address space
- OS binary image (kernel code, data) is mapped into **the virtual address space of every process at addresses not used by process** (high VA)
- Why is this done? Easy to jump to OS code during a trap



# A subtle point

- OS is not a separate process with its own address space
- Instead, OS code is part of the address space of every process
- A process sees OS as part of its code (e.g., like a library)
- During trap, process jumps to high virtual addresses and executes OS code



# OS is part of address space of every process

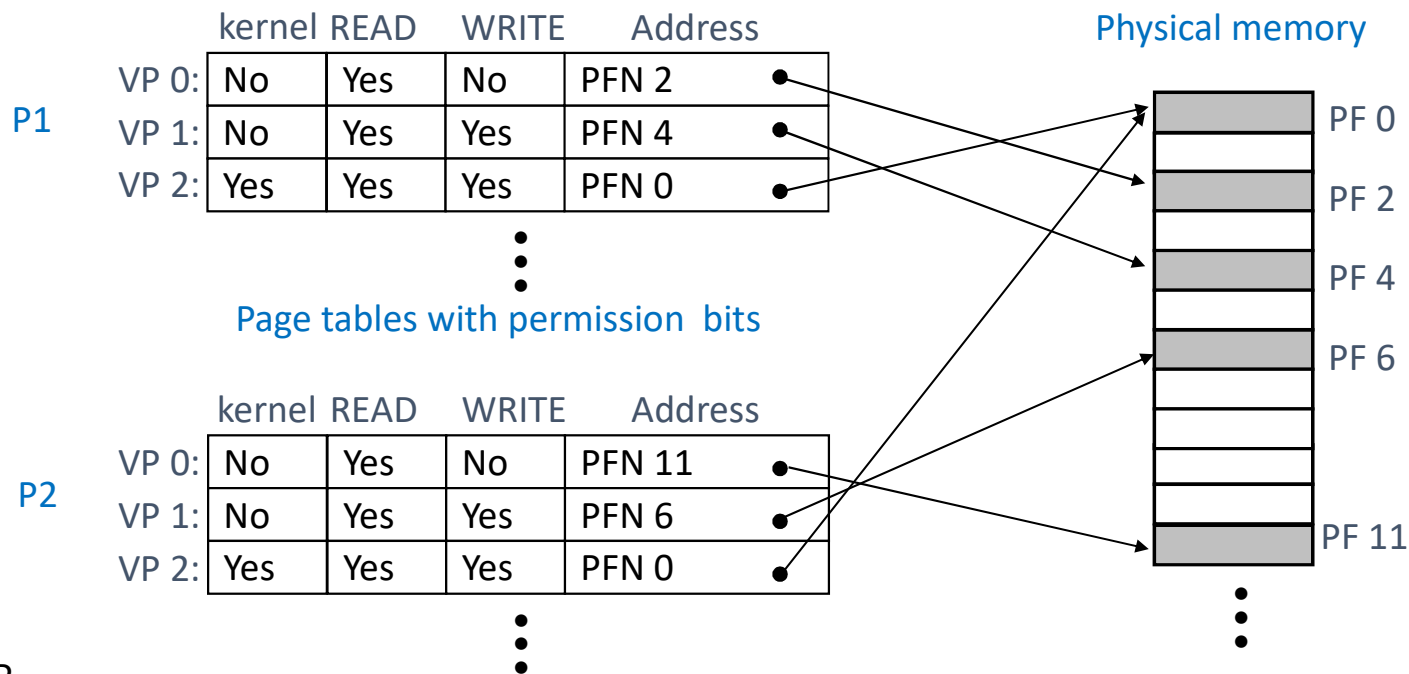
- OS code/data assigned virtual addresses
  - Compiler ensures high virtual addresses not used by user code
- OS virtual addresses are mapped to physical addresses of OS via page table entries of every process
- There is only one copy of OS code/data in RAM
  - Loaded into RAM at low physical addresses during system bootup
- Page tables of all processes have mappings to same OS physical addresses
  - Same high virtual addresses map to same physical addresses of OS code

# Page-level isolation and security

- How is OS code/data protected from illegal access by user?
- Page table has permissions for every memory page
  - Whether read/write or read-only (code pages are read-only)
  - Whether page can be accessed in user mode or kernel mode
- Page table mappings for OS code are protected to allow access only when CPU is in kernel mode
  - CPU in user mode cannot access high virtual addresses of OS code
  - CPU in kernel mode (after trap instruction) can access OS code/data
- MMU traps to OS if any violation detected during memory access, ensures user programs can only access memory they are permitted to access

# Example: page-level protection using page tables

- Example: process P1 and P2 each have one read-only page, one read-write page, and one page with OS code accessible in kernel mode



# Memory management in xv6

- 32-bit OS, so  $2^{32}=4\text{GB}$  virtual address space for every process
- 4KB pages, so 32 bit VA = 20 bit page number + 12 bit offset
- Each PTE has 20 bit physical frame number, and some flags
  - PTE\_P indicates if page is valid/present (if not set, access will cause page fault)
  - PTE\_W indicates if writeable (if not set, only reading is permitted)
  - PTE\_U indicates if user page (if not set, only kernel can access the page)
- Address translation: use page number (top 20 bits of virtual address) to index into page table, find physical frame number, add 12-bit offset

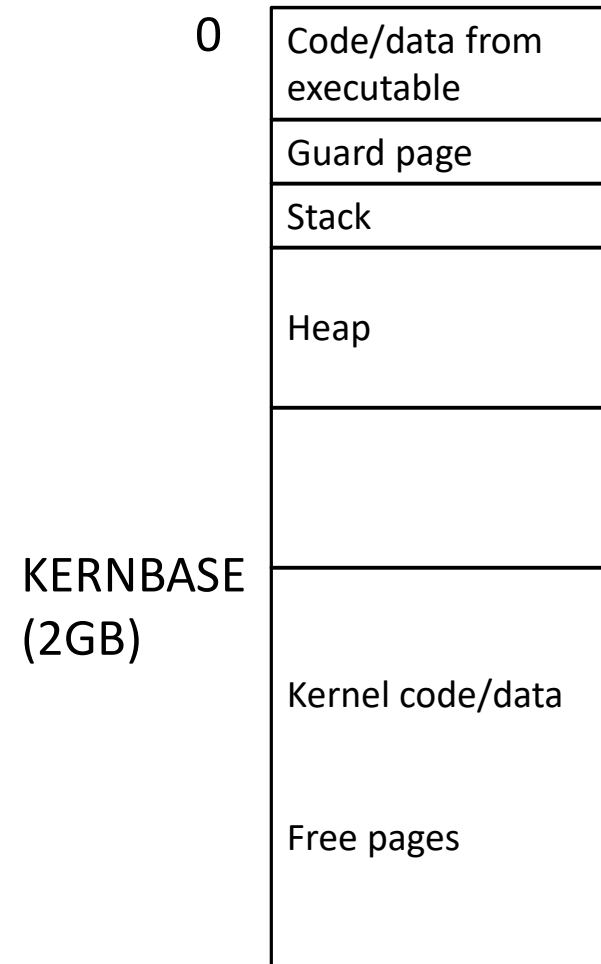
# Two level page table in xv6

- xv6 has two-level page table
  - 1024 “inner” page table pages, each with 1024 PTEs
  - Outer page directory stores PTE-like references to 1024 inner page table pages
  - Physical address of outer page directory is stored in CPU’s cr3 register, used by MMU during address translation
- 32 bit virtual address = 10 bits index into page directory, next 10 bits index into inner page table, last 12 bits are offset within page
  - PFN from PTE + offset = physical address

```
0773 // A virtual address 'la' has a three-part structure as follows:
0774 //
0775 // +-----10-----+-----10-----+-----12-----+
0776 // | Page Directory |   Page Table   | Offset within Page |
0777 // |      Index      |      Index      |                   |
0778 // +-----+-----+-----+
0779 // \--- PDX(va) ---/ \--- PTX(va) ---/
0780
```

# Virtual address space in xv6

- Virtual address space [ 0, 4GB]
- Physical address space [0, PHYSTOP] where PHYSTOP is max physical memory that can be used
- Virtual address space contains
  - Low virtual addresses: user code/data, guard page, stack, expandable heap
  - High virtual address starting at KERNBASE (2GB): kernel code/data, free pages that OS assigns to user processes, memory reserved for I/O devices, ...



# Page table mappings

- Page table contains two sets of PTEs
- User entries: low VA to PA used to process for code, data, stack, heap
- Kernel entries: high VA to PA containing OS code/data/free pages
  - $[KERNBASE, KERNBASE + PHYSTOP]$  mapped to  $[0, PHYSTOP]$
- Kernel page table entries identical across all processes

