


Lecture 29: Locking in xv6

Mythili Vutukuru

IIT Bombay

<https://www.cse.iitb.ac.in/~mythili/os/>

Why locking in xv6?



- No threads in xv6, so no two user programs can access same userspace memory image
 - No need for userspace locks like pthreads mutex
- However, scope for concurrency in xv6 kernel
 - Two processes in kernel mode in different CPU cores can access same kernel data structures
 - When a process is running in kernel mode, another trap occurs, and the trap handler can access data that was being accessed by previous kernel code
- Solution: spinlocks used to protect critical sections
 - Limit concurrent access to kernel data structures that can result in race conditions
- xv6 also has a sleeping lock (built on spinlock, not discussed)

P1 → kernel → intr

Spinlocks in xv6

- Acquiring lock: uses xchg x86 atomic instruction (test and set)
 - Atomically set lock variable to 1 and returns previous value
 - If previous value is 0, it means free lock has been acquired, success!
 - If previous value is 1, it means lock is held by someone, continue to spin in a busy while loop till success
- Releasing lock: set lock variable to 0
- Must disable interrupts on CPU core before spinning for lock
 - Interrupts disabled only on this CPU core to prevent another trap handler running and requesting same lock, leading to deadlock
 - OK for process on another core to spin for same lock, as the process on this core will release it
 - Disable interrupts before starting to spin (otherwise, vulnerable window after lock acquired and before interrupts disabled)

```
1500 // Mutual exclusion lock.
1501 struct spinlock {
1502     uint locked;           // Is the lock held?
1503
1504     // For debugging:
1505     char *name;            // Name of lock.
1506     struct cpu *cpu;       // The cpu holding the lock.
1507     uint pcs[10];          // The call stack (an array of program counters;
1508                           // that locked the lock.
1509 };
```

```
1573 void
1574 acquire(struct spinlock *lk)
1575 {
1576     pushcli(); // disable interrupts to avoid deadlock.
1577     if(holding(lk))
1578         panic("acquire");
1579
1580     // The xchg is atomic.
1581     while(xchg(&lk->locked, 1) != 0)
1582         ;
1583
1584     // Tell the C compiler and the processor to not move loads or stores
1585     // past this point, to ensure that the critical section's memory
1586     // references happen after the lock is acquired.
1587     __sync_synchronize();
1588
1589     // Record info about lock acquisition for debugging.
1590     lk->cpu = mycpu();
1591     getcallerpcs(&lk, lk->pcs);
1592 }
```

Handwritten notes:

- = 0 ✓ success*
- spin*
- C1 P1 → L → ~~intr.~~*
- C2 P2 → (L) spinning ✓*

Disabling interrupts

L1 L2

- Must disable interrupts on CPU core before beginning to spin for spinlock
- Interrupts stay disabled until lock is released
- What if multiple spinlocks are acquired?
 - Interrupts must stay disabled until all locks are released
- Disabling/enabling interrupts:
 - pushcli disables interrupts on first lock acquire, increments count for future locks
 - popcli decrements count, reenables interrupts only when all locks released and count is zero

```
1662 // Pushcli/popcli are like cli/sti except that they are matched:
1663 // it takes two popcli to undo two pushcli. Also, if interrupts
1664 // are off, then pushcli, popcli leaves them off.
1665
1666 void
1667 pushcli(void)
1668 {
1669     int eflags;
1670
1671     eflags = readeflags();
1672     cli();
1673     if(mycpu()->ncli == 0)
1674         mycpu()->intena = eflags & FL_IF;
1675     mycpu()->ncli += 1;
1676 }
1677
1678 void
1679 popcli(void)
1680 {
1681     if(readeflags() & FL_IF)
1682         panic("popcli - interruptible");
1683     if(--mycpu()->ncli < 0)
1684         panic("popcli");
1685     if(mycpu()->ncli == 0 && mycpu()->intena)
1686         sti();
1687 }
```

ptable.lock (1)

```
2409 struct {  
2410     struct spinlock lock;  
2411     struct proc proc[NPROC];  
2412 } ptable;
```

- The process table protected by a lock, any access to ptable must be done with ptable.lock held
- Normally, a process in kernel mode acquires ptable.lock, changes ptable, releases lock
 - Example: when allocproc allocates new struct proc
- But during context switch from process P1 to P2, ptable structure is being changed all through context switch, so when to release lock?
 - P1 acquires lock, switches to scheduler, switches to P2, P2 releases lock

If the process is context switched out from P1 to P2, and lock held, then the destination process releases the lock held by the initial process.

ptable.lock (2)

- Every function that calls sched() to give up CPU will do so with ptable.lock held. Which functions invoke sched()?
 - Yield, when a process gives up CPU due to timer interrupt
 - Sleep, when process wishes to block
 - Exit, when process terminates
- Every function that swtch switches to will release ptable.lock. What functions does swtch return to?
 - Yield, when switching in a process that is resuming after yielding is done
 - Sleep, when switching in a process that is waking up after sleep
 - Forkret for newly created processes
- Purpose of forkret: to release ptable.lock after context switch, before returning from trap to userspace

```
2826 // Give up the CPU for one scheduling round.
2827 void
2828 yield(void)
2829 {
2830     acquire(&ptable.lock);
2831     myproc()->state = RUNNABLE;
2832     sched();
2833     release(&ptable.lock);
2834 }
```

P1 → P2

```
2852 void
2853 forkret(void)
2854 {
2855     static int first = 1;
2856     // Still holding ptable.lock from scheduler.
2857     release(&ptable.lock);
2858
2859     if (first) {
2860         // Some initialization functions must be run in the context
2861         // of a regular process (e.g., they call sleep), and thus cannot
2862         // be run from main().
2863         first = 0;
2864         iinit(ROOTDEV);
2865         initlog(ROOTDEV);
2866     }
```

trapret

That's why we don't directly return from the trapret but instead we take the path of forkret to release the lock and then return.

ptable.lock (3)

- Scheduler goes into loop with lock held
- Switch to P1, P1 switches back to scheduler with lock held, scheduler switches to P2, P2 releases lock
- Periodically, end of looping over all processes, releases lock temporarily
 - What if no runnable process found due to interrupts being disabled?
Release lock, enable interrupts, allow processes to become runnable.

```
2757 void
2758 scheduler(void)
2759 {
2760     struct proc *p;
2761     struct cpu *c = mycpu();
2762     c->proc = 0;
2763
2764     for(;;){
2765         // Enable interrupts on this processor.
2766         sti();
2767
2768         // Loop over process table looking for process to run.
2769         acquire(&ptable.lock);
2770         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2771             if(p->state != RUNNABLE)
2772                 continue;
2773
2774             // Switch to chosen process. It is the process's job
2775             // to release ptable.lock and then reacquire it
2776             // before jumping back to us.
2777             c->proc = p;
2778             switchvm(p);
2779             p->state = RUNNING;
2780
2781             swtch(&(c->scheduler), p->context);
2782             switchkvm();
2783
2784             // Process is done running for now.
2785             // It should have changed its p->state before coming back.
2786             c->proc = 0;
2787         }
2788         release(&ptable.lock);
2789     }
2790 }
2791 }
```

The diagram illustrates the execution flow of the scheduler. A central node 'S' represents the scheduler, and two nodes 'P1' and 'P2' represent processes. Red arrows show the sequence of execution: the scheduler (S) switches to process P1, P1 switches back to the scheduler (S), the scheduler (S) switches to process P2, and P2 switches back to the scheduler (S). A large red arrow points from the 'release(&ptable.lock);' line in the code to the 'S' node, indicating the point where the scheduler releases the lock and resumes execution.

Summary

- Spinlocks in xv6 based on xchg atomic instruction
- Processes in kernel mode hold spinlock when accessing shared data structures, disabling interrupts on that core while lock is held
- Special ptable.lock held across context switch